

Lab Handbook

Machine and Human Intelligence group, University of Helsinki

2025-09-02

Table of contents

Introduction	4
I Python tools	5
1 How to publish Python packages	6
1.1 Publishing on PyPI	6
Other notes	9
1.2 Uploading to conda-forge	9
1.3 Updating the package version	11
II HPC resources	12
2 Introduction	13
2.1 CSC	13
2.2 HPC tutorials	13
3 Useful commands	14
3.1 Squeue	14
4 HPC Workflow	15
4.1 Scope	15
4.2 Initial Preparation	15
4.3 Running Jobs	16
Interactive Job	17
Batch Job	17
4.4 Monitoring and Logs	18
Monitoring Your Job	18
Inspecting Logs	19
4.5 References	19

III Research topics	20
5 Resources for Gaussian processes and Bayesian optimization	21
Gaussian Processes	21
Bayesian Optimization	21
Slack Channel	21
6 Seminars	22
6.1 Ongoing	22
6.2 Interesting Channels	22
7 Machine Learning Coffee Seminar organizing guide	23
7.1 Prerequisites	23
Zoom Meeting Ownership Transfer	23
7.2 Tasks	24
Pre-Seminar Tasks	24
7.3 Seminar Day Checklist	24
Equipment Setup	24
Zoom Meeting Management	25
Attendance Tracking	25
7.4 Post-Seminar Tasks	25
YouTube Upload Instructions	26
IV Other info	27
8 Contributing Guide	28
8.1 Steps to Contribute	28
8.2 Additional Resources	29
Bibliography	30

Introduction

Lab handbook of the [Machine and Human Intelligence group](#) at the University of Helsinki.

Part I

Python tools

1 How to publish Python packages

We describe here the steps and tips for publishing Python packages on PyPI and conda-forge.

1.1 Publishing on PyPI

Publishing a package on PyPI is generally simpler than publish on conda-forge, and it's easiest to use the PyPI package as the source for a conda-forge recipe anyway, so it's best to start here. More details can be found [here](#).

1. Make sure your package dependencies are all correctly specified, e.g. by creating a new Conda environment and running `pip install -e .` from an up-to-date source and ensuring that all tests pass. Dependencies should be specified with a lower bound, rather than pinned to a specific version, for maximum compatibility.
 - There are a few different ways to specify project dependencies, but the most modern is with a [pyproject.toml](#) file in the root of the project directory. This is what we have used for PyVBMC and GPyReg, this guide may need to be adapted if you are using a different method. From what I understand, it is also good practice to leave a “stub” in setup.py for backwards compatibility, as in [here](#).
2. Make sure `pyproject.toml` contains the appropriate lines. These are taken from PyVBMC, so they may not all be required for your project, but this is a starting point. You may also already have a `pyproject.toml` with other information in it, such as configurations for the [Black formatter](#). That's fine, those sections can remain.

```
# pyproject.toml
#
[project]
name = "PyVBMC" # Naming a project is required
dynamic = ["version"] # use git tags for version, via setuptools_scm
# If you don't want to use setuptools_scm, you can specify the version
# manually:
# version = "v1.0.0"
description = "Variational Bayesian Monte Carlo in Python."
readme = "README.md"
```

```

license = { file = "LICENSE" }
dependencies = [
    "cma >= 3.1.0",
    "corner >= 2.2.1",
    # ...
]
requires-python = ">=3.9"
#
[tool.setuptools]
include-package-data = true
# We want to include some files which are not in the source directory,
# such as the Jupyter Notebooks in /examples.
# See https://setuptools.pypa.io/en/latest/userguide/datafiles.html
# for more info.
packages = ["pyvbm", "pyvbm.examples"]
package-dir = {"pyvbm.examples" = "examples"}
#
[tool.setuptools.package-data]
# Specify the extensions to include as data:
"pyvbm.examples" = ["*.ipynb"]
# Including files which *are* in the source directory but are not
# .py files can also be accomplished with a MANIFEST.in file. See
# e.g. https://github.com/acerbilab/pyvbm/blob/main/MANIFEST.in
#
[project.optional-dependencies]
dev = [
    "myst_nb >= 0.13.1",
    "numpydoc >= 1.2.1",
    # ...
] # These dependencies are for developers. They will only be installed
# with `pip install pyvbm[dev]` (or `pip install .[dev]`, locally).
# You can specify other sets of optional dependencies similarly.
#
[build-system]
# Required for using setuptools.scm, which automatically assigns the
# version number based on Git tags.
requires = [
    "setuptools >= 45",
    "setuptools_scm[toml] >= 6.2",
]
build-backend = "setuptools.build_meta"

```

- `setuptools_scm` is a tool which extracts versioning info from Git tags, so that you don't have to manually specify package versions.

3. Tag the current commit as a release version with

```
git tag vX.Y.Z
```

See [here](#) for some details about semantic versioning.

4. Install Python build tools. It's probably a good idea to create a new environment for the whole build and packaging process.

```
conda create -n build-env # optional, but recommended
conda activate build-env # if you created build-env, otherwise make sure you have activated y
pip install setuptools-scm
pip install build
pip install twine
```

5. Build your package:

```
python -m build
```

(from the project directory)

6. This should create a directory `dist/` with `.whl` and `.tar.gz` files matching your package name and tagged version. Inspect these files and make sure that they contain the contents you expect. By default this should be the source directory corresponding to your project's name, e.g. `/pyvbm`. If you are missing files that should be there, then see the link above (<https://setuptools.pypa.io/en/latest/userguide/datafiles.html>) regarding `package-data` and `MANIFEST.in`.

7. You can also run `twine check dist/*` to ensure that the package name and description will render properly on PyPI.

8. If everything looks correct, test out the build by creating a new environment and running `pip install dist/*.whl` to install the packaged version and test it out. For example, you could run the tests with `pytest --pyargs your-package` (do this from somewhere *outside* the project directory, to ensure that `pytest` is finding the version you just installed, and not the local tests). You could also open a Python REPL and just ensure that your package imports. You may need to open a new shell before the installation can be found on your path.

```
conda create -n test-package-release
conda activate test-package-release
pip install dist/*.whl
cd ~ # Test the build package outside the working directory.
```



```
pytest --pyargs your-package
# Ensure later to come back to the project working directory
```

9. If everything checks out and you are ready to upload, first head to <https://test.pypi.org/> and register an account if you don't already have one. Then you can test uploading your package by running `twine upload --repository testpypi dist/*` from the project directory. Your package should then be visible under your account on the test repository, and you can make sure that the description and other info are correct. You can also attempt to install it with

```
python3 -m pip install --index-url https://test.pypi.org/simple/ --extra-index-url https://p
```

The `--extra-index-url` tells `pip` to also look at the regular PyPI repository, which is important if your package has dependencies which are not on the test repo (this is likely the case).

10. If everything looks good on the test repository, you can run `twine upload dist/*` to upload to the official repository (you will need an account there as well, separate from your test account). Be aware that while a version of a package can be deleted from PyPI (so that it is no longer available), that same version number can never be re-uploaded. So it pays to double-check.

11. Once uploaded, you should be able to run `pip install your-package!`

Other notes

1. You can add another user as an owner or maintainer of a project through your account at <https://pypi.org/>
2. Optionally, you may also want to list the package version on GitHub. You can do this by pushing the version tag(s) to the origin with `git push --tags`, then selecting the tag from <https://github.com/account-name/repo-name/tags> and clicking “Create release”. You can describe the release, and optionally upload the built `.whl` and `.tar.gz` files.

1.2 Uploading to conda-forge

Uploading to `conda-forge` is slightly more involved. Detailed instructions can be found [here](#), but here is a summary of steps:

1. Fork the `conda-forge` [staged-recipes](#) repository from GitHub, and clone the fork locally. Checkout a new branch, e.g. `pyvbmc-recipe`, then `cd` into the `staged-recipes/recipes/` directory.

2. Install `grayskull` with `pip install grayskull` (or `conda install -c conda-forge grayskull`). `Grayskull` is a utility which will help generate the appropriate metadata for your package.
3. If your package is already on PyPI, you can run `grayskull pypi --strict-conda-forge your-package` to generate a `conda-forge` recipe for your package. It should create the file `staged-recipes/recipes/your-package/meta.yaml`.
4. Everything should be filled in automatically (though it's good to double-check), with the exception of

```
|about:
|  home: https://acerbilab.github.io/pyvbm/
|
|
|extra:
|  recipe-maintainers:
|    - AddYourGitHubIdHere
```

(ignore the vertical bars, Colab won't let me include leading spaces without them). You can also add other people as recipe maintainers, with their permission. It just means they will receive automated PRs and other updates regarding the project from `conda-forge`, and possibly answer any questions that pop up. 5. Optionally, you can add documentation and development URLs to the `about:` section, e.g.

```
|about:
|  home: https://acerbilab.github.io/pyvbm/
|  description: |
|    PyVBM is a Python implementation of the Variational Bayesian Monte Carlo (VBM) algorithm
|  doc_url: https://acerbilab.github.io/pyvbm/
|  dev_url: https://github.com/acerbilab/pyvbm/
|  ...
```

6. Optionally, you can include commands in the `test` section of the `meta.yaml` recipe, which will be automatically run when `conda-forge` updates the package. The `pyvbm` test suite is quite expensive to run, so I elected to just include the default basic tests here, which just ensure that the package can be imported.
7. Once you think everything is correct (see a checklist [here](#)), you can commit the changes to your new branch, push them to GitHub, and then open a PR to merge your fork to the original `staged-recipes` repo. Fill out the template checklist which appears when you draft the PR.
8. Once you've opened the PR, the `conda-forge` automation will check your recipe and attempt to install the package. Correct any errors that occur, and comment on GitHub with `@conda-forge-admin, please restart ci` to re-run the automated checks.

9. Once the automated checks all pass, you can ping a member of the `conda-forge` team to review and approve the PR with `@conda-forge-admin`, please ping `conda-forge/help-python`.
10. After the PR is merged, it will take a few hours (and possibly up to 24) for the package to become available on the Conda servers. After that `conda install --channel=conda-forge your-package` should work!
11. A repo will be created at `conda-forge/your-package-feedstock`, and you and any other recipe maintainers will be added to it. This is where automated PRs regarding your package will be issued.

1.3 Updating the package version

Fortunately, this part is simple.

- **Updating the version on PyPI:**

Re-run steps 3-5 from the section [Publishing on PyPI](#) above, and upload the new `.whl` and `.tar.gz` files to PyPI with `twine`.

- **Updating the version on conda-forge:**

After the package becomes available on PyPI, `conda-forge` should automatically find the new version (this will take a few hours). After that, the `conda-forge` process will automatically issue a PR to the feedstock repo, which a recipe maintainer can approve in order to update the `conda-forge` version. After approval, the binary on the `conda-forge` cloud/repo will be automatically updated within 24 hours.

Part II

HPC resources

2 Introduction

Whenever doing research, High Performance Computing (HPC) access and usage is fundamental, you will rarely be running things only on a local machine. The main HPC resources we can use are [CSC](#) and [UH's Turso](#).

2.1 CSC

The CSC is in charge of Finland's national computational resources. Many storage and computing programmes exist here, but the main ones for running ML experiments are Lumi, Mahti, and Puhti.

- Mahti and Puhti are considered the same project resource. When you get approved for a project, you can use either to run. The Mahti partition contains 24 GPU nodes, each equipped with four NVIDIA A100 GPUs. The Puhti partition includes 80 GPU nodes, each featuring four NVIDIA Volta V100 GPUs.
- Lumi is separate from Mahti and Puhti and has a separate project approval process. It contains 2,978 nodes, each equipped with four AMD Instinct MI250X GPUs, totaling 11,912 GPUs, considered one of Europe's largest supercomputers

2.2 HPC tutorials

This is a [tutorial on HPC usage](#) held in June 2021 by Aalto University, very useful. The videos are [here](#).

[HPC user guide for UH](#). They have [support sessions](#) organized every day if you have any questions or need any help with HPC usage.

Slack channel: Join and check out the [#hpc](#) channel on the lab Slack.

3 Useful commands

3.1 Squeue

Check the status of submitted slurm jobs.

```
squeue -u $USER --format="%.18i %.9P %.8j %.8u %.8T %.10M %.20S %.9l %.6D %R %.10m"
```

For convenience, you could set the default by writing the following lines in your `.bashrc` or `.bash_profile` and then simply call `squeue`.

```
alias squeue='squeue -u $USER'  
export SQUEUE_FORMAT="%.18i %.9P %.8j %.8u %.8T %.10M %.20S %.9l %.6D %R %.10m"
```

See the [squeue manual](#) for more options.

4 HPC Workflow

This guide explains how to use High-Performance Computing (HPC) resources available at the University of Helsinki (UH). Currently, we have access to several HPC systems:

1. [Turso](#)
2. [Puhti and Mahti](#)
3. [LUMI](#)

In most cases, LUMI is the preferred choice for training large-scale models, despite its AMD GPU architecture. Puhti and Mahti are scheduled to be replaced by [Rouhu](#) in 2026. Rouhu has not yet been launched at the time of writing this guide.

4.1 Scope

This workflow focuses on using **LUMI**. However, the steps described here can be easily adapted to other HPC systems. Typically, only a few settings—such as the partition name, account name, and GPU or CPU resource specifications—need to be modified when switching clusters. About 80% of the workflow remains the same, assuming the system uses [SLURM](#) as its scheduler.

4.2 Initial Preparation

Before you begin, ensure that:

- You have active [CSC accounts](#).
- You are connected to the university network (or connected via [VPN](#)).

Then, prepare your working environment as follows:

1. **Log in to the cluster**

For example, to log in to LUMI:

```
ssh <your_csc_username>@lumi.csc.fi
```

This connects you to the *login node*. **Do not** run resource-intensive jobs on the login node—it is shared and has limited resources.

2. **Navigate to your working directory**

On LUMI, switch to your [scratch space](#) and create a project folder:

```
cd /scratch/<project>
mkdir -p <your_username>/<your_project>
cd <your_username>/<your_project>
```

Do not store large project files (such as data and code) in your home directory, as its storage space is limited. Also note that data stored in the scratch space are temporary and may be deleted automatically—keep that in mind when organizing your work.

3. **Upload your data and code**

Transfer datasets, scripts, and source code into your project directory using `scp`, `rsync`, or `git clone`.

4. **Set up your software environment**

There are several ways to manage Python environments on HPC systems. This guide does not cover those in detail, but for LUMI-specific guidance, refer to the official documentation: [Python on LUMI](#)

With your environment ready, you can start submitting and managing SLURM jobs.

4.3 Running Jobs

When submitting a job to LUMI, you must specify the target partition. A list of available partitions is available in the LUMI documentation: [LUMI Partitions](#). Check the availability and type of resources offered by each partition—for example, whether your job requires GPUs or can run on CPU-only nodes.

There are two main ways to run jobs:

1. **Interactive jobs** – useful for debugging or interactive exploration.
2. **Batch jobs** – for longer, unattended runs using a SLURM batch script.

Always start with a small test job to verify that your script runs correctly. This helps catch issues early and avoid wasting large-scale resources.

Interactive Job

Interactive sessions allow you to log in to a compute node with an allocated shell, ideal for testing and debugging.

```
salloc --nodes=1 --account=<project> --partition=<partition> --time=00:30:00
```

Explanation of flags:

- `--nodes=1`: allocates one node
- `--account=<project>`: your LUMI project name
- `--partition=<partition>`: e.g., `small`, `standard`, or `standard-g`
- `--time=00:30:00`: time limit in HH:MM:SS

After allocation, you'll get an interactive shell on the compute node. Note that the required flags and their formats might differ across HPC systems. For more details, see LUMI's [Interactive Jobs Guide](#).

Batch Job

For longer, automated runs, use a SLURM batch script. Here is a minimal example:

Create a script `run_job.sh`:

```
#!/bin/bash

# SBATCH directives
#SBATCH --account=<project>           # Project/account to charge
#SBATCH --job-name=<your_job_name>    # Descriptive job name
#SBATCH --partition=standard-g       # Queue/partition (e.g., GPU-enabled)
#SBATCH --nodes=1                    # Number of nodes
#SBATCH --ntasks-per-node=1          # MPI tasks per node
#SBATCH --cpus-per-task=16           # CPU cores per task
#SBATCH --gres=gpu:1                 # Number of GPUs
#SBATCH --mem=32GB                    # RAM requested
#SBATCH --time=05:00:00              # Max wall time (HH:MM:SS)
#SBATCH --output=slurm_log/%x_%A.out # Log path: jobname_jobID.out

# Environment setup
export PATH="/scratch/<project>/<env_name>/bin:$PATH"
cd $SLURM_SUBMIT_DIR
```

```
export PYTHONPATH=$PWD

#      Execute
srun python -m your_python_script.py
```

Submit the job using the `sbatch` command, which sends your batch script to the SLURM scheduler:

```
sbatch run_job.sh
```

Wait for the job to enter the queue and start. Running multiple jobs in parallel is possible; refer to the official documentation for advanced array jobs: [LUMI Throughput](#).

4.4 Monitoring and Logs

Monitoring Your Job

After you submit a job with `sbatch`, SLURM assigns it a unique job ID. You can use this ID to monitor the job's status and resource usage.

To check the status of all your jobs:

```
squeue --user=$USER
```

To check the status of a specific job:

```
squeue -j <jobID>
```

To automatically refresh the job status every few seconds, use the `watch` command. For example:

```
watch squeue --user=$USER
```

This keeps re-running the `squeue` command and updates the display every 2 seconds by default.

To monitor resource usage of a running or completed job:

```
seff <jobID>
```

If your job is underutilizing requested resources, adjust future jobs accordingly to save compute credits.

You can also run `slurm history` to view a list of your past jobs, including their IDs, states, and runtimes. This can be useful for tracking completed jobs or troubleshooting issues.

Inspecting Logs

Logs are saved in the path specified in your script (under `slurm_log/` if you use above script). These log files usually include standard output (stdout) and standard error (stderr) streams from your job:

```
slurm_log/<your_job_name>_<jobID>.out
```

Check logs after each run—they contain important information for debugging and performance tuning.

4.5 References

- <https://scicomp.aalto.fi/triton/tut/intro/#example-project>

Part III

Research topics

5 Resources for Gaussian processes and Bayesian optimization

Gaussian Processes

- **Gortler et al. (2019)**. A Visual Exploration of Gaussian Processes. [Distill](#)
- **Videos and tutorials** from the Gaussian Process Summer School: [GPSS 2021 Program](#)
 - Start with the first video, “Intro to GPs”, and the sessions on “Kernel Design” and “Representation Learning with GPs” (other videos can be explored later as time permits).
- **Tutorials/workshops (Jupyter notebooks)**: [GPSS 2021 Labs](#)
- **The GP Bible** by Rasmussen and Williams (2006): [Gaussian Processes for Machine Learning](#)

Bayesian Optimization

- **Exploring Bayesian Optimization** (Agnihotri and Batra 2020). [Distill](#)
- **Frazier (2018)**. A Tutorial on Bayesian Optimization. [arXiv](#)
- **Video**: “Introduction to Bayesian Optimization” from the Gaussian Process Summer School, Day 3: [GPSS 2021 Day 3](#)
- **The Bayesian Optimization Book** by Garnett (2023): [BayesOptBook](#)

Slack Channel

- Join and check out the [#gaussian-processes](#) channel on the lab Slack.

6 Seminars

6.1 Ongoing

- [Machine Learning Coffee seminars](#)
- [Aalto Seminar on Advances in Probabilistic Machine Learning \[APML\]](#)
- [Seminar on Gaussian Processes, Spatiotemporal Modeling, and Decision-making Systems](#)

6.2 Interesting Channels

- [Secondmind Labs Research seminars](#)

7 Machine Learning Coffee Seminar organizing guide

7.1 Prerequisites

To organize the ML seminar, ensure you have the following access and resources:

- **University of Helsinki user account**
- **Key to Computer Science Department** to access room A318

Before beginning your tasks, complete these access requests:

1. **FCAI Squarespace and YouTube channel access:** Contact Kaisa Pekkala (kaisa.pekkala@aalto.fi).
2. **FCAI speakers and participants documents:** Request access from Luigi.
3. **Zoom Meeting Ownership:** Request ownership transfer from the previous organizer.

Zoom Meeting Ownership Transfer

1. Sign in to Zoom with your University of Helsinki account.
 2. Go to **Settings** → **Meetings**.
 3. Scroll to the bottom and select **Schedule Privilege** → **Edit and Add**.
 4. Enter the previous owner's university email and select their name.
 5. Ask the previous owner to transfer room ownership to you:
 - They should go to Zoom, locate **MLCS Kumpula** under **Meetings**, choose **Edit and Schedule For**, select your name, and save.
 - [More information on Zoom ownership transfer](#).
-

7.2 Tasks

Pre-Seminar Tasks

1. **Create Event on Squarespace:**
 - Duplicate an existing event if preferred, but update **Options** → **Event URL** to reflect the new event.
 2. **Send Advertising Email:**
 - Confirm with the ML seminar team (mlseminar@hiit.fi) that the announcement has been sent.
 3. **Reserve Microphone:**
 - Reserve a **Blackstorm Scout USB microphone** from IT support (helpdesk@helsinki.fi).
 4. **Print Attendance List**
-

7.3 Seminar Day Checklist

Equipment Setup

1. **Get Equipment:**
 - Pick up the **Blackstorm Scout microphone** (USB) from IT support (the microphone is stored in room A216, has black and purple box).
 - Retrieve the **Owl camera** from meeting room A318 (stored in a white box on the side table by the door).
2. **Classroom Setup:** Setting up the classroom for the hybrid seminar (in-person and Zoom attendees) can be tricky so ensure to test in advance for the first time. There are two main setups to consider:
 - **Using your own laptop without the classroom computer:**
 - Connect the microphone and Owl camera to your laptop. You could test in advance to ensure they work with your laptop.
 - Position the Owl camera on the table in front of the podium.
 - Start the Zoom meeting from your laptop

- Let the speaker join the Zoom meeting from their laptop, muting the speaker’s laptop microphone and turning off the camera. The speaker should use the microphone connected to your laptop.
- Connect the speaker’s laptop to the projector and share their screen.
- If the projector doesn’t work with the speaker’s laptop, you can connect your laptop for sharing the speaker’s screen, which is less ideal since it will have some Zoom toolbar visible on the screen.
- **Start recording** in Zoom.
- **Using the classroom computer:**
 - Connect the microphone and Owl camera to the classroom computer. Some classroom computers might not work well, so you might need to use your laptop instead. See the previous setup for details.
 - Start the projector.
 - Start the Zoom meeting from your laptop.
 - Let the speaker join the Zoom meeting from their laptop.
 - Log into Zoom from the classroom computer (you’ll need the university account for login) and share the meeting to the projector screen.
 - **Start recording** in Zoom

Zoom Meeting Management

1. Record the session on your laptop.
2. Monitor Zoom chat and read questions aloud if needed.
3. Note the number of Zoom attendees.

Attendance Tracking

- Circulate the attendance list and a pen once the talk begins.
 - Make a note of the Zoom attendance count.
-

7.4 Post-Seminar Tasks

1. **Return Equipment:**
 - Return the microphone and Owl camera to their respective storage locations.
2. **Handle Leftover Food:**
 - Move any leftover food to the coffee room on the second floor of the CS department.

3. **Record Attendance:**

- Update the attendance in the FCAI participants document.

4. **Edit and Upload Video:**

- Use QuickTime or a similar tool for video editing.
- Upload the edited video to YouTube.

YouTube Upload Instructions

1. **Title:** Use the title from the event information.
2. **Description:** Include the abstract and speaker details from the event information.
3. **Playlist:** Add the video to the **Machine Learning Coffee Seminar** playlist.
4. **Audience:** Set the audience to **not made for kids**.

Part IV

Other info

8 Contributing Guide

Here are the detailed steps to ensure a smooth workflow. **It's simply editing the .qmd file and making pull requests on GitHub.** So don't be scared .

8.1 Steps to Contribute

1. Fork the Repository

Visit the [handbook repository](#) and click on “Fork” to create a copy of the repository under your GitHub account.

2. Clone the Repository

Clone your fork to your local machine:

```
git clone https://github.com/YOUR_USERNAME/handbook.git
cd handbook
```

3. Create a New Branch

Before making changes, create a new branch to keep your work organized:

```
git checkout -b your-branch-name
```

4. Add or Modify Content

- Write or edit a (q)markdown file as needed.
- If you're adding a new chapter, create a new .qmd file, then include it under the `book: chapters:` field in `__quarto.yml`. For example:

```
book:
  chapters:
    - index.qmd
    - part: Python tools
    - your-new-chapter.qmd
```

5. Preview Locally

To see your changes before pushing, build and preview the book locally. You could use for example VSCode as the editor and Quarto has a VSCode plugin to help preview the book website locally, see instructions [here](#).

Or make sure you have Quarto installed, then run in bash terminal:

```
quarto preview
```

This will open a live preview in your browser where you can verify your changes.

6. Commit and Push Changes

After making sure your changes look good, commit and push them to your fork:

```
git add .  
git commit -m "Brief description of your changes"  
git push origin your-branch-name
```

7. Make a Pull Request

Go to the original [handbook repository](#) and open a Pull Request from your branch. Provide a brief summary of your changes and mention any specific details reviewers should be aware of.

8.2 Additional Resources

For more advanced features and customization options, please refer to the [Quarto documentation](#).

Bibliography

- Agnihotri, Apoorv, and Nipun Batra. 2020. “Exploring Bayesian Optimization.” *Distill*. <https://doi.org/10.23915/distill.00026>.
- Garnett, Roman. 2023. *Bayesian Optimization*. Cambridge University Press.
- Rasmussen, Carl Edward, and Christopher KI Williams. 2006. *Gaussian Processes for Machine Learning*. Vol. 2. 3. MIT press Cambridge, MA.